

# Undirected graphical models

*and other unnormalised (aka energy-based) models*

Pierre-Alexandre Mattei



**3iA** Côte d'Azur  
Institut interdisciplinaire  
d'intelligence artificielle

## Menu for this lecture

---

1. Motivation for undirected models
2. Undirected graphical models, general theory
3. Examples of undirected models
4. Learning in undirected graphical models and energy-based models

# 1

Motivation(s) for undirected  
models

## What are directed graphical models?

---

- Let  $G = (V, E)$  be a DAG whose nodes are denoted  $V = \{1, \dots, d\}$
- Let  $X = (X_1, \dots, X_d)$  be a random variable with density  $p(x) = p(x_1, \dots, x_d)$ .
- Every node of the graph corresponds to a random variable

# What are directed graphical models?

---

- Let  $G = (V, E)$  be a DAG whose nodes are denoted  $V = \{1, \dots, d\}$
- Let  $X = (X_1, \dots, X_d)$  be a random variable with density  $p(x) = p(x_1, \dots, x_d)$ .
- Every node of the graph corresponds to a random variable

**Definition** (Directed graphical model). We say that  $p$  **factorises in**  $G$  (denoted  $p \in \mathcal{L}(G)$ ) when, for all  $x$ ,

$$p(x) = \prod_{i=1}^d p(x_i | x_{\text{pa}_i}),$$

where, for all node  $i$ ,  $\text{pa}_i$  denotes the set of parents of node  $i$ .

# What are directed graphical models?

- Let  $G = (V, E)$  be a DAG whose nodes are denoted  $V = \{1, \dots, d\}$
- Let  $X = (X_1, \dots, X_d)$  be a random variable with density  $p(x) = p(x_1, \dots, x_d)$ .
- Every node of the graph corresponds to a random variable

**Definition** (Directed graphical model). We say that  $p$  *factorises in*  $G$  (denoted  $p \in \mathcal{L}(G)$ ) when, for all  $x$ ,

$$p(x) = \prod_{i=1}^d p(x_i | x_{\text{pa}_i}),$$

← the likelihood (or bounds at least) can often be computed exactly


where, for all node  $i$ ,  $\text{pa}_i$  denotes the set of parents of node  $i$ .

## Motivating example: image segmentation with a neural net (eg Unet)

We have couples  $(x_i, y_i)_{i \leq n}$




1000 | X



## Motivating example: image segmentation with a neural net (eg Unet)

We have couples  $(x_i, y_i)_{i \leq n}$

- $x_i \in \mathbb{R}^{n_{\text{pixels}}}$
- $y_i \in \{1, \dots, C\}^{n_{\text{pixels}}}$






## Motivating example: image segmentation with a neural net (eg Unet)

We have couples  $(x_i, y_i)_{i \leq n}$

- $x_i \in \mathbb{R}^{n_{\text{pixels}}}$
- $y_i \in \{1, \dots, C\}^{n_{\text{pixels}}}$

number of possible segmentation classes  
(e.g. bicycle, cyclist, horse, background...)



## Simple examples: image segmentation with a neural net (eg Unet)

$$p_{\beta}(x_1, \dots, x_n, y_1 \dots y_n) = \prod_{i=1}^n p(x_i) p_{\beta}(y_i | x_i)$$

We want to use a Unet that takes as input an image and gives as output one probability vector per pixel.

## Simple examples: image segmentation with a neural net (eg Unet)

$$p_{\beta}(x_1, \dots, x_n, y_1 \dots y_n) = \prod_{i=1}^n p(x_i) p_{\beta}(y_i | x_i)$$



We want to use a Unet that takes as input an image and gives as output one probability vector per pixel.

$$p_{\beta}(y_i | x_i) = \prod_{j=1}^{n_{\text{pixels}}} \mathcal{M}(y_{ij} | [\text{Unet}_{\beta}(x_i)]_j)$$


## Simple examples: image segmentation with a neural net (eg Unet)

$$p_{\beta}(x_1, \dots, x_n, y_1 \dots y_n) = \prod_{i=1}^n p(x_i) p_{\beta}(y_i | x_i)$$



We want to use a Unet that takes as input an image and gives as output one probability vector per pixel.

$$p_{\beta}(y_i | x_i) = \prod_{j=1}^{n_{\text{pixels}}} \mathcal{M}(y_{ij} | [\text{Unet}_{\beta}(x_i)]_j)$$




## Simple examples: image segmentation with a neural net (eg Unet)

$$p_{\beta}(x_1, \dots, x_n, y_1 \dots y_n) = \prod_{i=1}^n p(x_i) p_{\beta}(y_i | x_i)$$



We want to use a Unet that takes as input an image and gives as output one probability vector per pixel.

$$p_{\beta}(y_i | x_i) = \prod_{j=1}^{n_{\text{pixels}}} \mathcal{M}(y_{ij} | [\text{Unet}_{\beta}(x_i)]_j)$$



**Q.** Do you see issues with this model?


## Simple examples: image segmentation with a neural net (eg Unet)

$$p_{\beta}(x_1, \dots, x_n, y_1 \dots y_n) = \prod_{i=1}^n p(x_i) p_{\beta}(y_i | x_i)$$



We want to use a Unet that takes as input an image and gives as output one probability vector per pixel.




$$p_{\beta}(y_i | x_i) = \prod_{j=1}^{n_{\text{pixels}}} \mathcal{M}(y_{ij} | [\text{Unet}_{\beta}(x_i)]_j)$$



**Q.** Do you see issues with this model?

**The pixels maybe shouldn't be independent given the image...**

How do we fix this model ?



## Some issues with DAGs

---

- As we saw in the initial lectures, a graph is a DAG if and only if it can be topologically ordered
- Hence, DAGs are mostly well-suited for cases where there is a **natural ordering over modalities**



## Some issues with DAGs

---

- As we saw in the initial lectures, a graph is a DAG if and only if it can be topologically ordered
- Hence, DAGs are mostly well-suited for cases where there is a **natural ordering over modalities**
- We saw previously that, **when there are two natural orderings (eg in VAEs or diffusion models), we could just use two DAGs**, but often that's not the case

## Some issues with DAGs

---

- As we saw in the initial lectures, a graph is a DAG if and only if it can be topologically ordered
- Hence, DAGs are mostly well-suited for cases where there is a **natural ordering over modalities**
- We saw previously that, **when there are two natural orderings (eg in VAEs or diffusion models), we could just use two DAGs**, but often that's not the case
- Even then, two DAGs means two models, which is not always satisfying.

## Another motivating example


- This is a nice model if we're just interested in pure supervised learning
- But we can't sample images from this
- We can't compute  $p(x)$ , handle missing data
- It is tempting to reverse the arrow, that would allow us to do all that...
- Of course, we could do two DAGs, but can we do a single model?

**Will Grathwohl**  
University of Toronto & Vector Institute  
Google Research  
wgrathwohl@cs.toronto.edu

**Kuan-Chieh Wang\* & Jörn-Henrik Jacobsen\***  
University of Toronto & Vector Institute  
wangkual@cs.toronto.edu  
j.jacobsen@vectorinstitute.ai

**David Duvenaud**  
University of Toronto & Vector Institute  
duvenaud@cs.toronto.edu

**Kevin Swersky & Mohammad Norouzi**  
Google Research  
{kswersky, mnorouzi}@google.com



$$p_{\beta}(y_i|x_i) = \text{Categorical}(y_i|\text{CNN}_{\beta}(x))$$

YOUR CLASSIFIER IS SECRETLY AN ENERGY BASED  
MODEL AND YOU SHOULD TREAT IT LIKE ONE

## Another motivating example


- This is a nice model if we're just interested in pure supervised learning

**Will Grathwohl**  
University of Toronto & Vector Institute  
Google Research  
wgrathwohl@cs.toronto.edu

**Kuan-Chieh Wang\* & Jörn-Henrik Jacobsen\***  
University of Toronto & Vector Institute  
wangkua1@cs.toronto.edu  
j.jacobsen@vectorinstitute.ai

**David Duvenaud**  
University of Toronto & Vector Institute  
duvenaud@cs.toronto.edu


**Kevin Swersky & Mohammad Norouzi**  
Google Research  
{kswersky, mnorouzi}@google.com



$$p_{\beta}(y_i|x_i) = \text{Categorical}(y_i|\text{CNN}_{\beta}(x))$$

## Another motivating example


- This is a nice model if we're just interested in pure supervised learning
- But we can't sample images from this
- We can't compute  $p(x)$ , handle missing data



$$p_{\beta}(y_i|x_i) = \text{Categorical}(y_i|\text{CNN}_{\beta}(x))$$

## Another motivating example

- This is a nice model if we're just interested in pure supervised learning
- But we can't sample images from this
- We can't compute  $p(x)$ , handle missing data
- It is tempting to reverse the arrow, that would allow us to do all that...



$$p_{\beta}(y_i|x_i) = \text{Categorical}(y_i|\text{CNN}_{\beta}(x))$$

# YOUR CLASSIFIER IS SECRETLY AN ENERGY BASED MODEL AND YOU SHOULD TREAT IT LIKE ONE

## Another motivating example


- This is a nice model if we're just interested in pure supervised learning
- But we can't sample images from this
- We can't compute  $p(x)$ , handle missing data
- It is tempting to reverse the arrow, that would allow us to do all that...

**Will Grathwohl**  
University of Toronto & Vector Institute  
Google Research  
wgrathwohl@cs.toronto.edu

**Kuan-Chieh Wang\* & Jörn-Henrik Jacobsen\***  
University of Toronto & Vector Institute  
wangkual@cs.toronto.edu  
j.jacobsen@vectorinstitute.ai

**David Duvenaud**  
University of Toronto & Vector Institute  
duvenaud@cs.toronto.edu

**Kevin Swersky & Mohammad Norouzi**  
Google Research  
{kswersky, mnorouzi}@google.com



$$p_{\beta}(y_i|x_i) = \text{Categorical}(y_i|\text{CNN}_{\beta}(x))$$

- Of course, we could do two DAGs, **but can we do a single model?**

# 2

Undirected graphical models  
**(aka Markov random fields)**



## Image segmentation as a motivation for undirected models


---

- We would like to model some correlation between neighbourhoods of pixels after the image exits the convnet.
- Unfortunately, **it is a bit arbitrary to look at pixels as nodes of a DAG**: indeed, there is no obvious notion of ordering on the pixels of an image!
- It would make sense to **create a graph without order**, for which nearby pixels are connected. This is exactly the sort of thing we can do with undirected models!

# What's a clique?

---

**Definition.** *A totally connected subset of vertices is called a **clique**.*



## What's an undirected graphical model?

---

**Definition.** Let  $G = (V, E)$  be an undirected graph. We denote by  $\mathcal{C}$  the set of all cliques of  $G$ . We say that a **probability distribution**  $p$  **factorizes in  $G$**  and write  $p \in \mathcal{L}(G)$  if  $p(x)$  is of the form:

$$p(x) = \frac{1}{Z} \prod_{C \in \mathcal{C}} \psi_C(x_C) \quad \text{with} \quad \psi_C \geq 0, C \in \mathcal{C} \quad \text{and} \quad Z = \int \prod_{C \in \mathcal{C}} \psi_C(x_C) dx.$$

The functions  $\psi_C$  are called **factors** or **potentials**.

$\subset + \infty$

## What's an undirected graphical model?

---

**Definition.** Let  $G = (V, E)$  be an undirected graph. We denote by  $\mathcal{C}$  the set of all cliques of  $G$ . We say that a **probability distribution**  $p$  **factorizes in  $G$**  and write  $p \in \mathcal{L}(G)$  if  $p(x)$  is of the form:

$$p(x) = \frac{1}{Z} \prod_{C \in \mathcal{C}} \psi_C(x_C) \quad \text{with} \quad \psi_C \geq 0, C \in \mathcal{C} \quad \text{and} \quad Z = \int \prod_{C \in \mathcal{C}} \psi_C(x_C) dx.$$

The functions  $\psi_C$  are called **factors** or **potentials**.

- Contrarily to directed models, **the factors are not necessarily probability densities!**
- Why can we be sure that, in spite of this,  $p(x)$  will still be a proper density?

## What's an undirected graphical model?

---

**Definition.** Let  $G = (V, E)$  be an undirected graph. We denote by  $\mathcal{C}$  the set of all cliques of  $G$ . We say that a **probability distribution**  $p$  **factorizes in  $G$**  and write  $p \in \mathcal{L}(G)$  if  $p(x)$  is of the form:


$$p(x) = \frac{1}{Z} \prod_{C \in \mathcal{C}} \psi_C(x_C) \quad \text{with} \quad \psi_C \geq 0, C \in \mathcal{C} \quad \text{and} \quad Z = \int \prod_{C \in \mathcal{C}} \psi_C(x_C) dx.$$

The functions  $\psi_C$  are called **factors** or **potentials**.

- Contrarily to directed models, **the factors are not necessarily probability densities!**
- Why can we be sure that, in spite of this,  $p(x)$  will still be a proper density?
  - Because of the division by  $Z$ , that ensures that  $p(x)$  sums to one.
- Because of this normalisation, multiplying the potentials by constants will not change anything.

## The simplest example: the trivial graph


---



- What are all the possible cliques of this graph?

## The simplest example: the trivial graph

---




- What are all the possible cliques of this graph?
  - Just the three individual nodes!
- Therefore, the distributions that factorise in this trivial graph will be of the form

$$p(x) = \frac{1}{Z} \psi_1(x_1) \psi_2(x_2) \psi_3(x_3)$$

## The simplest example: the trivial graph

---



- What are all the possible cliques of this graph?
  - Just the three individual nodes!
- Therefore, the distributions that factorise in this trivial graph will be of the form

$$p(x) = \frac{1}{Z} \psi_1(x_1) \psi_2(x_2) \psi_3(x_3)$$


- A direct consequence is that  $\forall i \in \{1, 2, 3\}, p(x_i) = \frac{1}{\int \psi_i(x) dx} \psi_i(x_i)$ , and therefore

$$p(x) = p(x_1)p(x_2)p(x_3)$$

- Like in the directed case, **the trivial graph corresponds to independence.**



## The other end of the spectrum: the complete graph




- Again, what are the possible cliques?
  - All the **individual nodes**, but also **the pairs**, and also **the full graph**!



1000 | X

## The other end of the spectrum: the complete graph




- Again, what are the possible cliques?
  - All the **individual nodes**, but also **the pairs**, and also **the full graph**!

$$p(x) = \frac{1}{Z} \psi_1(x_1) \psi_2(x_2) \psi_3(x_3) \psi_{12}(x_1, x_2) \psi_{23}(x_2, x_3) \psi_{13}(x_1, x_3) \psi_{123}(x_1, x_2, x_3)$$



## The other end of the spectrum: the complete graph



- Again, what are the possible cliques?

➤ All the **individual nodes**, but also **the pairs**, and also **the full graph**!

$$p(x) = \frac{1}{Z} \psi_1(x_1) \psi_2(x_2) \psi_3(x_3) \psi_{12}(x_1, x_2) \psi_{23}(x_2, x_3) \psi_{13}(x_1, x_3) \psi_{123}(x_1, x_2, x_3)$$

- By posing  $\tilde{\psi}_{123}(x_1, x_2, x_3) = \psi_1(x_1) \psi_2(x_2) \psi_3(x_3) \psi_{12}(x_1, x_2) \psi_{23}(x_2, x_3) \psi_{13}(x_1, x_3) \psi_{123}(x_1, x_2, x_3)$

this can be rewritten simply as  $p(x) = \frac{1}{Z} \tilde{\psi}_{123}(x_1, x_2, x_3)$ , so only looking at the largest clique

(called a **maximal clique**) will suffice.


- Like in the directed case, any distribution  $p$  will factorise in this complete graph by choosing

$$\tilde{\psi}_{123}(x_1, x_2, x_3) = p(x_1, x_2, x_3)$$

## Should we limit ourselves to maximal cliques

---


- **A maximal clique is a clique that cannot be included in a strictly larger clique.** Is this clique maximal?



## Should we limit ourselves to maximal cliques?

---

- **A maximal clique is a clique that cannot be included in a strictly larger clique.** Is this clique maximal?



- **Yes!** But all the cliques strictly included inside (e.g. the singletons or pairs) are not!

## Should we limit ourselves to maximal cliques?

---

- **A maximal clique is a clique that cannot be included in a strictly larger clique.**
- In the factorisation formula  $p(x) = \frac{1}{Z} \prod_{C \in \mathcal{C}} \psi_C(x_C)$ , we considered  $\mathcal{C}$  to be the set of all possible cliques. But like we did for the complete graph, **we could consider only one factor for each maximal clique** (by including inside all the factors of the smaller cliques).

## Should we limit ourselves to maximal cliques?

---

- **A maximal clique is a clique that cannot be included in a strictly larger clique.**
- In the factorisation formula  $p(x) = \frac{1}{Z} \prod_{C \in \mathcal{C}} \psi_C(x_C)$ , we considered  $\mathcal{C}$  to be the set of all possible cliques. But like we did for the complete graph, **we could consider only one factor for each maximal clique** (by including inside all the factors of the smaller cliques).
- This insight leads some authors **to change a bit the definition of undirected models, replacing  $\mathcal{C}$  by the set of all maximal cliques**. For instance, [Bishop] does this.
- The maximal clique definition reduces the number of factors, but make them harder to interpret. For instance, in the previous complete graph example, maybe  $\psi_{23}(x_2, x_3)$  had a specific interesting form, but considering a single factor obscures this. This is why we did not chose this version of the definition.

## Separation and undirected models

---

- For directed models, we had the nice concept of d-separation that helped us answer questions like «  $X_A \perp\!\!\!\perp X_B | X_C$ ? ». Do we have something similar for undirected models?
- Yes! Now it's just called **separation** (and the definition is much simpler).



## Separation and undirected models

- For directed models, we had the nice concept of d-separation that helped us answer questions like «  $X_A \perp\!\!\!\perp X_B | X_C$ ? ». Do we have something similar for undirected models?
- Yes! Now it's just called **separation** (and the definition is much simpler).
- **Separation recipe:** we consider all chains between any node in  $A$  and any node in  $B$ . Any of these chains is said to be **blocked** if it passes through  $C$ .




Fig. 8.27 from  
[Bishop]:  
C separates A and B.

## Properties of separation for undirected models

---

- Like in the directed case, the separation condition is sound and complete. The soundness constitutes a relatively famous theorem first proved by Hammersley and Clifford in 1971.

**Definition.** We say that  $p$  satisfies the **Global Markov property** w.r.t.  $G$  when, for all  $A, B, C$  disjoint subsets of  $V$ ,  
 $(A \text{ and } B \text{ are separated by } C) \Rightarrow (X_A \perp\!\!\!\perp X_B \mid X_C).$

**Theorem** (soundness of separation, Hammersley-Clifford). *If  $p > 0$  then*  
 $p \in \mathcal{L}(G) \iff p \text{ satisfies the global Markov property.}$

*Proof.* See e.g. [PGM, Sec. 4.3.1.1].

□

## Properties of separation for undirected models

---

- Like in the directed case, the separation condition is sound and complete. The completeness result is similar to the one for d-separation.

**Theorem** (completeness of separation). *If  $A$  and  $B$  are not separated by  $C$ , then there exist  $p \in \mathcal{L}(G)$  such that  $X_A \not\perp\!\!\!\perp X_B | X_C$ .*

*Proof.* See e.g. [PGM, Sec. 4.3.1.2].

□

## From directed to undirected ?

---

- The general properties of directed and undirected models are quite similar, are there ways to go from one to the other?



PROX



## From directed to undirected ?

---

- The general properties of directed and undirected models are quite similar, are there ways to go from one to the other?
- **Undirected seems more general than directed**, since the potentials are not required to be probability densities. So there should be a way to go from a directed graph to a more general undirected graph...
- We want to turn  $p(x) = \prod_{i=1}^d p(x_i | x_{\text{pa}_i})$ , into  $p(x) = \frac{1}{Z} \prod_{C \in \mathcal{C}} \psi_C(x_C)$ . How should we do this?

## From directed to undirected ?


---

- The general properties of directed and undirected models are quite similar, are there ways to go from one to the other?
- **Undirected seems more general than directed**, since the potentials are not required to be probability densities. So there should be a way to go from a directed graph to a more general undirected graph...
- We want to turn  $p(x) = \prod_{i=1}^d p(x_i | x_{\text{pa}_i})$ , into  $p(x) = \frac{1}{Z} \prod_{C \in \mathcal{C}} \psi_C(x_C)$ . How should we do this?
  1. We reverse all arrows to get an undirected graph. But that's not enough, since we need to turn each  $(x_i, x_{\text{pa}_i})$  into a clique.
  2. We « marry the parents » by drawing edges between them.

## Graph « moralisation »

- This recipe is called « moralising » a graph (because we marry the parents):
  1. We reverse all arrows to get an undirected graph. But that's not enough, since we need to turn each  $(x_i, x_{\text{pa}_i})$  into a clique.
  2. We « marry the parents » by drawing edges between them.

What is the moralised version of this graph?  
Fig. 8.33 in [Bishop].




## Graph « moralisation »

- This recipe is called « moralising » a graph (because we marry the parents):
  1. We reverse all arrows to get an undirected graph. But that's not enough, since we need to turn each  $(x_i, x_{\text{pa}_i})$  into a clique.
  2. We « marry the parents » by drawing edges between them.

**Let's just apply the recipe!**

Fig. 8.33 in [Bishop].





# 3

Some examples

# Conditional random fields for protein structure prediction

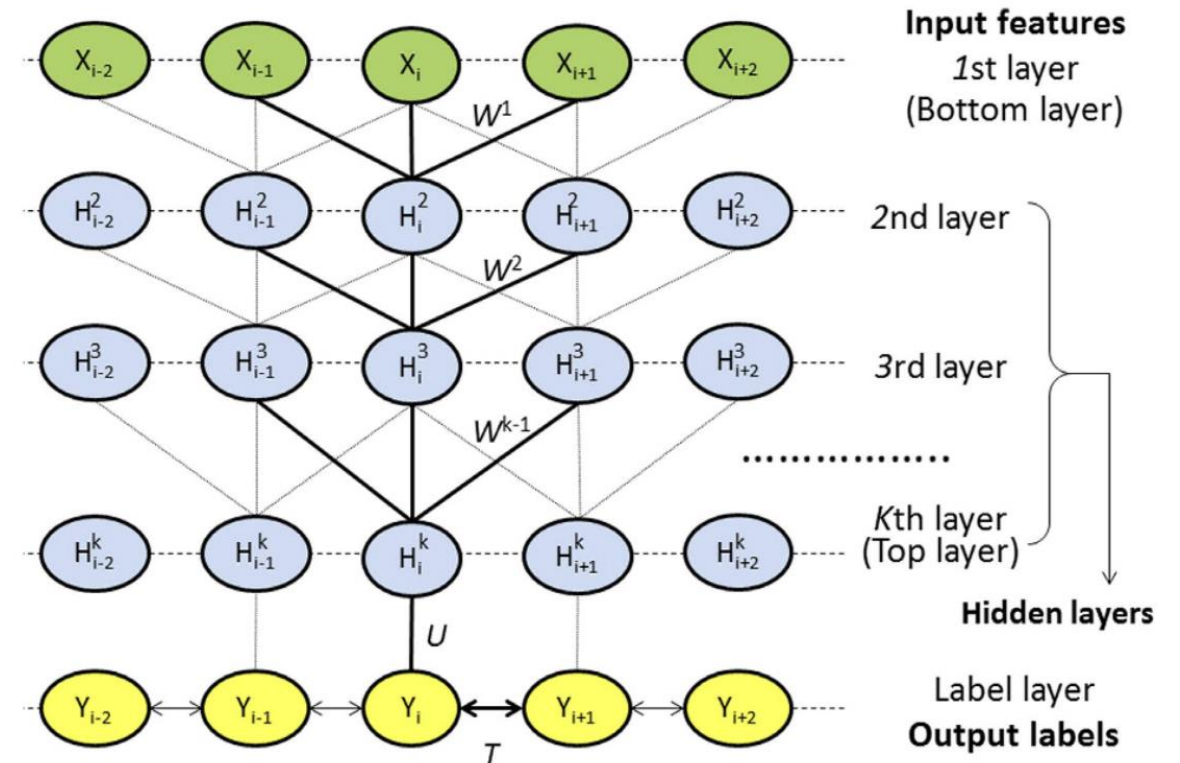
- The idea of using the outputs of a neural net as univariate potentials for images or sequence data is generally called a **conditional random field**

Article | [Open access](#) | Published: 11 January 2016

## Protein Secondary Structure Prediction Using Deep Convolutional Neural Fields

[Sheng Wang](#), [Jian Peng](#), [Jianzhu Ma](#) & [Jinbo Xu](#)

[Scientific Reports](#) **6**, Article number: 18962 (2016) | [Cite this article](#)



# Conditional random fields for image segmentation

---

- The idea of using the outputs of a neural net as univariate potentials for images or sequence data is generally called a **conditional random field**



## Conditional Random Fields as Recurrent Neural Networks

Shuai Zheng<sup>\*1</sup>, Sadeep Jayasumana<sup>\*1</sup>, Bernardino Romera-Paredes<sup>1</sup>, Vibhav Vineet<sup>†1,2</sup>, Zhizhong Su<sup>3</sup>,  
Dalong Du<sup>3</sup>, Chang Huang<sup>3</sup>, and Philip H. S. Torr<sup>1</sup>

<sup>1</sup>University of Oxford

<sup>2</sup>Stanford University

<sup>3</sup>Baidu Institute of Deep Learning

# Conditional random fields for image segmentation

---

- The idea of using the outputs of a neural net as univariate potentials for images or sequence data is generally called a **conditional random field**

10/10


## Conditional Random Fields as Recurrent Neural Networks

Shuai Zheng<sup>\*1</sup>, Sadeep Jayasumana<sup>\*1</sup>, Bernardino Romera-Paredes<sup>1</sup>, Vibhav Vineet<sup>†1,2</sup>, Zhizhong Su<sup>3</sup>,  
Dalong Du<sup>3</sup>, Chang Huang<sup>3</sup>, and Philip H. S. Torr<sup>1</sup>

<sup>1</sup>University of Oxford

<sup>2</sup>Stanford University

<sup>3</sup>Baidu Institute of Deep Learning



## Getting rid of a lot of constraints

- Exact likelihood models rely on complex constraints on the types of architectures used.
- **Can we get rid of these constraints, and consider super general deep generative models? Something like  $p_{\theta}(\mathbf{x}) = \text{NN}_{\theta}(x)$  with any neural net?**
- Does this make sense?

## Getting rid of a lot of constraints

- Exact likelihood models rely on complex constraints on the types of architectures used.
- **Can we get rid of these constraints, and consider super general deep generative models? Something like  $p_{\theta}(\mathbf{x}) = \text{NN}_{\theta}(x)$  with any neural net?**
- Does this make sense? Not really like this! We have at least **two very basic constraints that we can't get rid of...**
  - The density must **be non-negative**
  - The density must **sum to one**

## Getting rid of a lot of constraints


- Exact likelihood models rely on complex constraints on the types of architectures used.
- **Can we get rid of these constraints, and consider super general deep generative models? Something like  $p_{\theta}(\mathbf{x}) = \text{NN}_{\theta}(x)$  with any neural net?**
- Does this make sense? Not really like this! We have at least **two very basic constraints that we can't get rid of...**
  - The density must **be non-negative**
  - The density must **sum to one**
- This motivates the very general framework of **energy-based models**

$$p_{\theta}(\mathbf{x}) \propto \exp(f_{\theta}(x))$$

Any sort of neural net

## Getting rid of a lot of constraints

- EBMs are **undirected graphical models**, with density


$$p_{\theta}(\mathbf{x}) = \frac{\exp(f_{\theta}(x))}{Z_{\theta}}$$

Where  $Z_{\theta} = \int \exp(f_{\theta}(x))dx$  is the **normalising constant**.

- Aka **Boltzman distribution** in statistical physics, and  $E_{\theta}(x) = -f_{\theta}(x)$  is the **energy**
- When  $E_{\theta}$  is linear, this is called an **exponential family** in statistics



# 4

Training undirected models and  
other unnormalised models

## MLE for EBMs?

- EBMs are **undirected graphical models**, with density

$$p_{\theta}(\mathbf{x}) = \frac{\exp(f_{\theta}(x))}{Z_{\theta}}$$

Where  $Z_{\theta} = \int \exp(f_{\theta}(x))dx$  is the **normalising constant**.

- Computing the normalising constant seems daunting... Can we really do (approximate) maximum likelihood for these models?

## Getting rid of a lot of constraints

- EBMs are **undirected graphical models**, with density

$$p_{\theta}(\mathbf{x}) = \frac{\exp(f_{\theta}(x))}{Z_{\theta}}$$

Where  $Z_{\theta} = \int \exp(f_{\theta}(x))dx$  is the **normalising constant**.

- Computing the normalising constant seems daunting... Can we really do (approximate) maximum likelihood for these models?
- Remember that we do not need to actually compute the likelihood, but just to optimise it.
- So the only thing we really need are **gradients of the log-likelihood!**

## Getting rid of a lot of constraints

- EBMs are **undirected graphical models**, with density

$$p_{\theta}(\mathbf{x}) = \frac{\exp(f_{\theta}(x))}{Z_{\theta}}$$

Where  $Z_{\theta} = \int \exp(f_{\theta}(x))dx$  is the **normalising constant**.

- Computing the normalising constant seems daunting... Can we really do (approximate) maximum likelihood for these models?
- Remember that we do not need to actually compute the likelihood, but just to optimise it.
- So the only thing we really need are **gradients of the log-likelihood!**

## Computing log-likelihood gradients

- Let's try to compute our gradients!

Problematic term!!



$$\nabla \log p_\theta(\mathbf{x}) = \nabla f_\theta(x) - \nabla \log Z_\theta$$

- We'll now focus on the gradient of the log-normaliser,

$$\nabla \log Z_\theta = \frac{\nabla Z_\theta}{Z_\theta} = \frac{1}{Z_\theta} \int \nabla \exp(f_\theta(x)) dx$$

which gives

$$\nabla \log Z_\theta = \frac{1}{Z_\theta} \int \exp(f_\theta(x)) \nabla f_\theta(x) dx$$

*(Handwritten red annotations: a red circle around the fraction 1/Z\_\theta, a red arrow pointing from the integral to the text "= p(x)" written above it, and a red line connecting the circle to the integral.)*

and finally

$$\nabla \log Z_\theta = \int p(x) \nabla f_\theta(x) dx = \mathbb{E}_{x \sim p_\theta} [\nabla f_\theta(x)]$$

## Computing log-likelihood gradients

$$\nabla \log Z_{\theta} = \int p(x) \nabla f_{\theta}(x) dx = \mathbb{E}_{x \sim p_{\theta}} [\nabla f_{\theta}(x)]$$

gives us a way to estimate log-likelihood gradients, provided that we can sample from the model! It's often called **contrastive divergence**.

- It is possible to sample from the model using MCMC, for instance **Langevin Monte Carlo**.

## Aparté: Langevin Monte Carlo

- **Langevin Monte Carlo** provides a general way of sampling from distributions of the form

$$p(\mathbf{x}) \propto \exp(U(x))$$

- The idea is that this distribution is the stationary distribution of the Langevin diffusion equation  $X' = \nabla U(X) + \sqrt{2}W'$ , where  $W$  is a Brownian motion

- Solving the equation by discretising it gives

$$x^{k+1} = x^k + \tau \nabla U(x^k) + \sqrt{2\tau} \varepsilon_k$$

Stein score

where  $\tau$  is essentially a learning rate, and  $\varepsilon_k$  is standard normal.

- This allows to sample approximatively according to  $p(\mathbf{x}) \propto \exp(U(x))$

## How to train EBMs? Hyvärinen's score matching

---

- Key idea: the gradient of  $\log p_\theta(x)$  does not depend on the normalising constant!

$$\nabla_x \log p_\theta(x) = \nabla_x f_\theta(x)$$

- One way of having  $p_{\hat{\theta}} \approx p_{\text{data}}$  is to have  $\nabla_x \log p_{\hat{\theta}} \approx \nabla_x \log p_{\text{data}}$ , so Hyvärinen suggested to minimise

$$\mathbb{E} \left[ \left\| \nabla_x \log p_{\text{data}}(x) - \nabla_x \log p_\theta(x) \right\|^2 \right]$$



$$\text{Fisher score} = \nabla_{\theta} \log p_{\theta}(x)$$

## How to train EBM's? Hyvärinen's score matching

- Key idea: the gradient of  $\log p_{\theta}(x)$  does not depend on the normalising constant!

$$\nabla_x \log p_{\theta}(x) = \nabla_x f_{\theta}(x)$$

- One way of having  $p_{\hat{\theta}} \approx p_{\text{data}}$  is to have  $\nabla_x \log p_{\hat{\theta}} \approx \nabla_x \log p_{\text{data}}$ , so Hyvärinen suggested to minimise

$$\mathbb{E} \left[ \left\| \nabla_x \log p_{\text{data}}(x) - \nabla_x \log p_{\theta}(x) \right\|^2 \right]$$

Hyvärinen score or Stein score of the model

## How to train EBMs? Hyvärinen's score matching

- Key idea: the gradient of  $\log p_\theta(x)$  does not depend on the normalising constant!

$$\nabla_x \log p_\theta(x) = \nabla_x f_\theta(x)$$

- One way of having  $p_{\hat{\theta}} \approx p_{\text{data}}$  is to have  $\nabla_x \log p_{\hat{\theta}} \approx \nabla_x \log p_{\text{data}}$ , so Hyvärinen suggested to minimise

$$\mathbb{E} \left[ \left\| \nabla_x \log p_{\text{data}}(x) - \nabla_x \log p_\theta(x) \right\|^2 \right]$$

Hyvärinen score or Stein score of the data

## How to train EBMs? Hyvärinen's score matching

- Key idea: the gradient of  $\log p_\theta(x)$  does not depend on

$$\nabla_x \log p_\theta(x) = -\nabla_x \log p_{\text{data}}(x)$$

This looks difficult  
because we don't  
know

- One way of having  $p_{\hat{\theta}} \approx p_{\text{data}}$  is to have  $\nabla_x \log p_{\hat{\theta}} \approx \nabla_x \log p_{\text{data}}$ , which Hyvärinen suggested to minimise

$$\mathbb{E} [||\nabla_x \log p_{\text{data}}(x) - \nabla_x \log p_\theta(x)||^2]$$

Hyvärinen score or Stein score of the data

## How to train EBMs? Hyvärinen's score matching

- Key idea: the gradient of  $\log p_\theta(x)$  does not depend on

$$\nabla_x \log p_\theta(x) = -\nabla_x \log p_{\text{data}}(x)$$

This looks difficult  
because we don't  
know

- One way of having  $p_{\hat{\theta}} \approx p_{\text{data}}$  is to have  $\nabla_x \log p_{\hat{\theta}} \approx \nabla_x \log p_{\text{data}}$ , which Hyvärinen suggested to minimise

$$\mathbb{E} [||\nabla_x \log p_{\text{data}}(x) - \nabla_x \log p_\theta(x)||^2]$$

Hyvärinen score or Stein score of the data

## How to train EBM's? Hyvärinen's score matching

Stein

$$s_{\theta}(x) = \text{score}(\text{U-net})$$

- Hyvärinen managed to attack  $\mathbb{E} [||\nabla_x \log p_{\text{data}}(x) - \nabla_x \log p_{\theta}(x)||^2]$  using integration by parts, but this requires the Hessian of  $\log p_{\theta}(x)$ , it is pretty much impossible to use for deep models...
- Hyvärinen, Estimation of Non-Normalized Statistical Models by Score Matching, JMLR 2005
- Another solution was found by Pascal Vincent, who suggested to replace  $\nabla_x \log p_{\text{data}}$  by  $\nabla_x \log p_{\text{data},\sigma}$ , where  $p_{\text{data},\sigma}$  is a kernel density estimate of  $\log p_{\text{data}}$

## How to train EBMs? Vincent's denoising score matching

- Vincent, A Connection Between Score Matching and Denoising Autoencoders, Neural Computation, 2011
- Idea: replace  $\nabla_x \log p_{\text{data}}$  by  $\nabla_x \log p_{\text{data},\sigma}$ , where  $p_{\text{data},\sigma}$  is a kernel density estimate of  $\log p_{\text{data}}$

$$p_{\text{data},\sigma}(x) = \frac{1}{n} \sum_{i=1}^n \mathcal{N}(x|x_i, \sigma^2 I)$$

$$p_{\sigma_1} \dots p_{\sigma_T}$$